

Fixed-seed0 full8 accuracy: descriptive points only, no new CI or non-inferiority claim

Task metrics follow the frozen evaluator; MMLU is question-weighted within task.

llama8b: 8/24 policy/task points

qwen4b: 24/24 policy/task points

mistral7b: 24/24 policy/task points

